

TEXT MINING PROJECT: NATURAL LANGUAGE PROCESSING MODELS TO PREDICT CUSTOMER SENTIMENTS FROM AMAZON FINE FOOD REVIEW

[HTTPS://GITHUB.COM/YAY66/NLP-AMAZON-REVIEW/BLOB/MAIN/TEXT_MINING_PROJECT_NLP_AMAZON_REVIEWS.IPYNB](https://github.com/YAY66/NLP-AMAZON-REVIEW/blob/main/TEXT_MINING_PROJECT_NLP_AMAZON_REVIEWS.IPYNB)

ISM 6359 Data Mining



MARCH 11, 2023
YAYUAN ZHANG

Contents

| | |
|--|----|
| Abstract..... | 3 |
| 1. Background | 4 |
| 1.1 State Text Mining Tool | 4 |
| 1.2 Stakeholders | 4 |
| 1.3 Data & Source | 4 |
| 2. Data Analysis | 5 |
| 2.1 Workflow Analysis process | 5 |
| 2.2 Import libraries | 5 |
| 2.3 Data Understanding | 5 |
| 2.3.1 Basic information | 5 |
| 2.3.2 EDA..... | 7 |
| 2.3.3 Checking total unique values in each column..... | 9 |
| 2.3.4 Checking Missing values..... | 9 |
| 2.4 Data Pre-processing | 10 |
| 2.4.1 Adding new column for helpful%..... | 10 |
| 2.4.2 Adding new column for upvote metrics. | 10 |
| 2.4.3 Assigning different labels..... | 11 |
| 2.4.4 Creating Pivot Table and visualization for above table..... | 11 |
| 2.5 NLP – Converting text into vector & Data cleaning | 12 |
| 2.5.1 Removing the duplicates..... | 13 |
| 2.5.2 Lower casting | 13 |
| 2.5.3 Removing punctuation..... | 13 |
| 2.5.4 Removing of Stop-words..... | 14 |
| 2.5.5 Removing URLs..... | 14 |
| 2.6 Checking dataset before next step | 14 |
| 2.7 Binary classification label | 15 |
| 3. Build Model | 15 |
| 3.1 Fitting on train & test set | 15 |
| 3.2 Fetching Top 20 Positive & Top 20 Negative words | 16 |
| 3.3 Evaluate models by applying multiple NLP techniques & multiple ML algorithms | 17 |
| 3.3.1 Text classifier function | 17 |
| 3.3.2 Prediction function & Confusion Matrix | 19 |

| | |
|--|----|
| 3.3.3 Dummy classifier..... | 20 |
| 3.3.4 Summary | 21 |
| 3.4 Logistic regression model on TFIDF vectorizer | 21 |
| 3.5 Upvote Prediction | 22 |
| 3.6 Handling imbalanced data using RandomOverSampler | 23 |
| 3.7 Tweaking Parameters..... | 24 |
| 3W..... | 25 |
| What went well? | 25 |
| What did not go well? | 26 |
| What I would do differently next time?..... | 26 |

Abstract

The food industry is one of the most important industries in the global economy, which is also an industry with relatively strong competition. As one of the largest online retailers of food, Amazon must gain a deep understanding of customer preferences and sentiments to make informed decisions about its products.

This project focuses on developing natural language processing models to predict customer sentiments from Amazon Fine Food Reviews. Also, the project aims to help businesses gain insights into customer preferences and sentiments to make informed decisions. In the project, various data preprocessing techniques are detailedly described to prepare the data for analysis.

Additionally, various natural language processing techniques are employed to learn the relationships between the text data and the corresponding sentiment labels. The stakeholders can understand the strengths and weaknesses of their products and develop strategies and plan to improve their products' quality and customer satisfaction.

Keywords: Text Mining, Data Pre-Processing, Sentiment analysis, Amazon Fine Food Review, Natural Language Processing, Machine Learning.

1. Background

1.1 State Text Mining Tool

In this process, the Python programming language that acts on the PyCharm and Jupyter platforms is my text mining tool and environment.

1.2 Stakeholders

- ❑ Food industries: the food industries that sell their product on Amazon. Their businesses rely heavily on customer feedback to improve their products and increase customer satisfaction. Therefore, this project can benefit businesses in the food industry by providing insights into customer preferences and sentiment by predicting customer sentiment.
- ❑ Consumers: The consumers who purchase food products on Amazon also can benefit from this project. Businesses can improve product quality and increase customer satisfaction by analyzing customer sentiments, resulting in a better shopping experience for consumers.

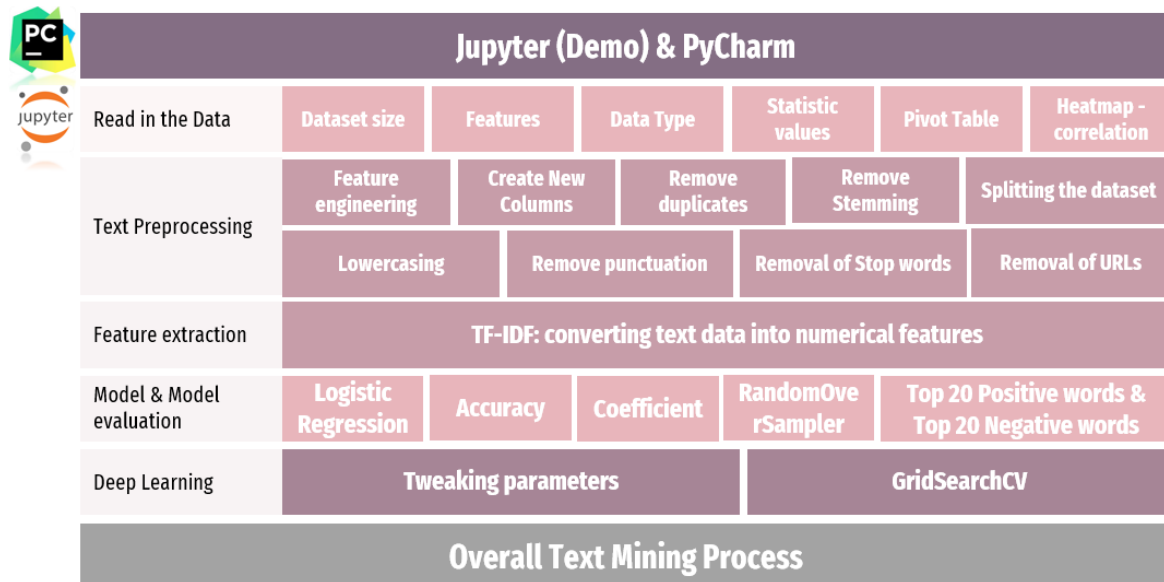
1.3 Data & Source

The dataset comes from Kaggle. The reviews were collected from October 1999 to October 2012. The dataset contains 568454 data points and 10 features. Each data point represents each customer reviews for various food products sold on Amazon, including the Product ID, Helpfulness Numerator, Helpfulness Denominator, Score, Summary, review text, etc.

The dataset is a fantastic resource for assessing consumer sentiments since it offers a thorough collection of client comments. The dataset's breadth and diversity also make it appropriate for training NLP models, facilitating the creation of precise sentiment analysis algorithms.

2. Data Analysis

2.1 Workflow Analysis process



2.2 Import libraries

First, I import the required libraries, including the science calculation, parts mainly for visualization, packages for pre-processing text data, model algorithms, ignoring some warnings, and so on.

2.3 Data Understanding

2.3.1 Basic information

The data has 568454 observations with 10 features. The target variable is the score. Other features are named: 'Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Time', 'Summary', and 'Text'.

```
# Checking the head of data and defining the target variable
```

```
df.head()
```

Checking the data frame size to learn what amount of data are dealing with

df.shape

In the dataset, 5 features are numeric fields, other 5 features are object fields.

Checking the datatypes for all features/columns & whether have empty cell or not

df.info()

```
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Id                                     568454 non-null  int64
1   ProductId                             568454 non-null  object
2   UserId                                 568454 non-null  object
3   ProfileName                            568438 non-null  object
4   HelpfulnessNumerator                   568454 non-null  int64
5   HelpfulnessDenominator                 568454 non-null  int64
6   Score                                  568454 non-null  int64
7   Time                                  568454 non-null  int64
8   Summary                                568427 non-null  object
9   Text                                   568454 non-null  object
dtypes: int64(5), object(5)
memory usage: 43.4+ MB
```

Next, I checked the distribution of the values of all the numeric fields. the describe function checks the distribution of the values of all the numeric fields, so it gets the counts, means, standard deviations, etc. for each feature.


Reading the data error description to check the spread of values across all numerical fields

data.describe().style.background_gradient(cmap='cividis')

| | Id | HelpfulnessNumerator | HelpfulnessDenominator | Score | Time | Helpful % |
|-------|---------------|----------------------|------------------------|---------------|-------------------|---------------|
| count | 568454.000000 | 568454.000000 | 568454.000000 | 568454.000000 | 568454.000000 | 568454.000000 |
| mean | 284227.500000 | 1.743817 | 2.228810 | 4.183199 | 1296256604.902420 | -0.067202 |
| std | 164098.679288 | 7.638513 | 8.289740 | 1.310436 | 48043312.332415 | 0.922175 |
| min | 1.000000 | 0.000000 | 0.000000 | 1.000000 | 939340800.000000 | -1.000000 |
| 25% | 142114.250000 | 0.000000 | 0.000000 | 4.000000 | 1271289600.000000 | -1.000000 |
| 50% | 284227.500000 | 0.000000 | 1.000000 | 5.000000 | 1311120000.000000 | 0.000000 |
| 75% | 426340.750000 | 2.000000 | 2.000000 | 5.000000 | 1332720000.000000 | 1.000000 |
| max | 568454.000000 | 886.000000 | 923.000000 | 5.000000 | 1351209600.000000 | 3.000000 |

2.3.2 EDA

Before any feature engineering, I prefer to quickly view the original data distribution by using data visualization to gain more insight and understand the relationship between each feature.




According to the score distribution bar chart above, a large number of reviews are 5 points, a total of 363122.

```
vc = df['Score'].value_counts().to_frame().reset_index()

fig = px.bar(x=vc["Score"], y=vc["index"],
             orientation='h', color=vc['index'], labels=True, color_continuous_scale="cividis")

fig.update_layout(title = "<b> Distribution Of The Score </b>", xaxis_title="", yaxis_title="<b>
Frequency </b> ", font_size=18, plot_bgcolor="#fff", showlegend=False)
```

Normally, if we fit highly correlated data in the model, it results in an overfitting problem. The good news is there is no highly correlated feature with the score in this data set.

```
plt.figure(figsize=(10, 10))
```

```
sns.heatmap(df.corr(), annot=True, cmap="cividis", annot_kws={"size":14})
```

2.3.3 Checking total unique values in each column.

I used the following coding to check the total unique values in each column.

```
dict={}

for i in df.columns:

    dict[i]=df[i].value_counts().shape[0]

pd.DataFrame(dict,index=['Unique']).T.style.background_gradient(cmap='cividis')
```

| | Unique |
|------------------------|--------|
| Id | 568454 |
| ProductId | 74258 |
| UserId | 256059 |
| ProfileName | 218416 |
| HelpfulnessNumerator | 231 |
| HelpfulnessDenominator | 234 |
| Score | 5 |
| Time | 3168 |
| Summary | 295742 |
| Text | 393579 |

2.3.4 Checking Missing values.

After checking the missing values and the percentage of missing values in the column by the following codes. Two features have missing values, Summary, and Profile Name. It seems like a small percentage. In this case, I made some optimizations to the code of past data mining projects, so that the number of missing values and its percentage can be directly displayed.

```
def missing_data(df):

    total = df.isnull().sum().sort_values(ascending = False)

    Percentage = (df.isnull().sum()/df.isnull().count()*100).sort_values(ascending = False)

    return pd.concat([total, Percentage], axis=1, keys=['Total', 'Percentage'])
```

```
missing_data(df)[missing_data(df)['Percentage']!=0].style.background_gradient(cmap='cividis')
```

| | Total | Percentage |
|-------------|-------|------------|
| Summary | 27 | 0.004750 |
| ProfileName | 16 | 0.002815 |

2.4 Data Pre-processing

Next, I will start doing data preprocessing to help improve the accuracy of the model.

2.4.1 Adding new column for helpful%

From the original dataset above I can see that there is a numerator and a denominator, so I will create a new column called "Helpful %" to add to the DataFrame (df for short) and calculate a value for each row representing that review "Useful ratio". Use the formula:

helpfulnessnumerator/helpfulnessdenominator. If the comment has not been rated by anyone, the value of this column is -1.

```
df['Helpful %'] = np.where(df['HelpfulnessDenominator'] > 0, df['HelpfulnessNumerator'] /
df['HelpfulnessDenominator'], -1)
```

2.4.2 Adding new column for upvote metrics.

Then, I cut the data into slices and make them into 7 bins, divide the column into 6 levels according to the percentage type as ['Empty' < '0-20%' < '20-40%' < '40-60%' < '60-80%' < '80-100 %'].

```
df['%upvote'] = pd.cut( df['Helpful %'] , bins = [-1, 0, 0.2, 0.4, 0.6, 0.8, 1.0], labels = ['Empty', '0-20%', '20-40%', '40-60%', '60-80%', '80-100%'])
```

```
df.head()
```

2.4.3 Assigning different labels

I assigned different labels to helpful% based on %upvote to its value, which counts data based on the score as a category and %upvote as a subcategory. In this step, I organized the table that shows only the ID column and reset the index.

```
df.groupby(['Score', '%upvote']).agg({'Id': 'count'})

df_s2=df.groupby(['Score', '%upvote']).agg({'Id': 'count'}).reset_index()
```

2.4.4 Creating Pivot Table and visualization for above table

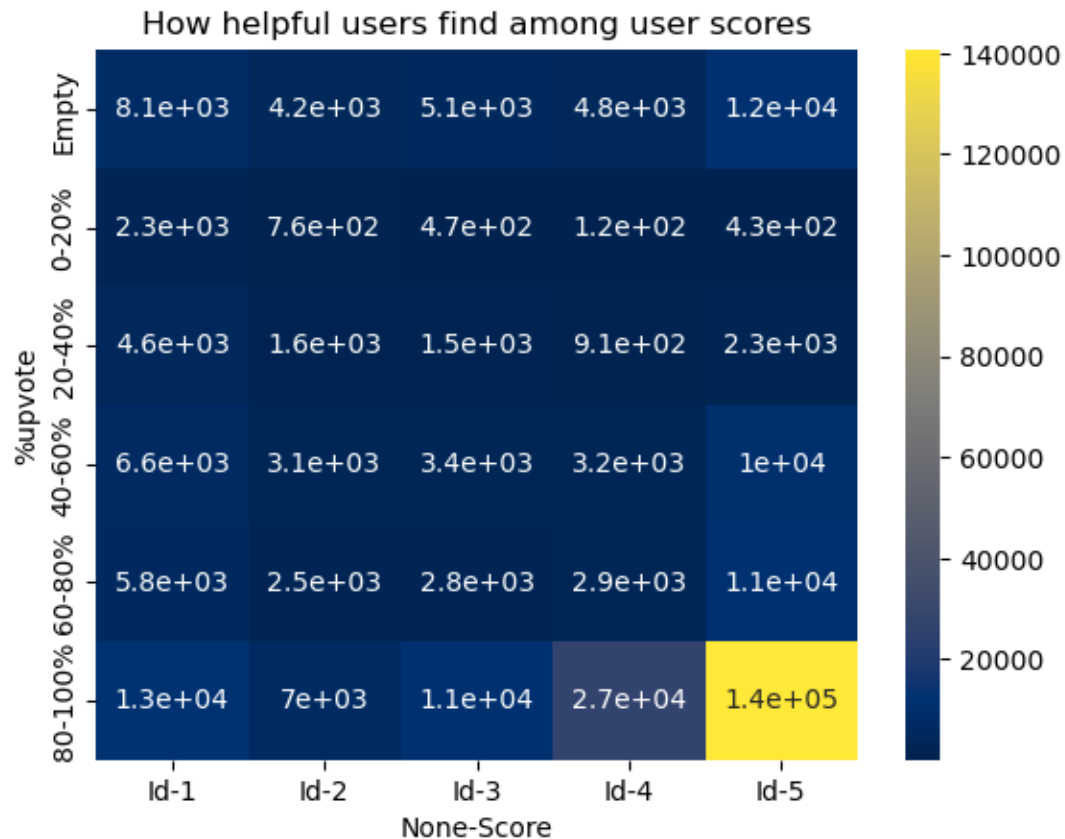
In this step, I created a PivotTable for better conclusions and a heatmap for it.

```
df_s2.pivot(index='%upvote',columns='Score')
```

| Score | 1 | 2 | 3 | 4 | 5 |
|---------|-------|------|-------|-------|--------|
| %upvote | | | | | |
| Empty | 8060 | 4234 | 5062 | 4780 | 11638 |
| 0-20% | 2338 | 762 | 474 | 116 | 432 |
| 20-40% | 4649 | 1618 | 1506 | 909 | 2275 |
| 40-60% | 6586 | 3051 | 3384 | 3185 | 10312 |
| 60-80% | 5838 | 2486 | 2754 | 2941 | 11060 |
| 80-100% | 12531 | 7014 | 11037 | 26707 | 140661 |

```
sns.heatmap(df_s2.pivot(index='%upvote',columns='Score'),annot=True,cmap = 'cividis')

plt.title('How helpful users find among user scores')
```



Based on the table above, it is seeing the key information:

- ☐ Reviews are skewed toward positive.
- ☐ More than half of the reviews are with zero votes.
- ☐ Lots of people rank 5-points for the review.

2.5 NLP – Converting text into vector & Data cleaning

Converting text into the vector is a key step in the NLP text mining project because computers cannot directly process text data. By converting text into vectors, we can use various machine learning algorithms to train models, classify, cluster, generate text, and many other tasks. Also, data often

contains some useless or redundant information, which interferes with the accuracy and efficacy of machine learning algorithms.

2.5.1 Removing the duplicates.

Regarding the data cleaning part, firstly I remove the duplicate data point if any by using the following coding.

```
df.duplicated().sum()

df.drop_duplicates(inplace=True)
```

2.5.2 Lower casting

Then, I convert all text to lowercase this avoids duplication or mistakes due to different cases.

```
df['Text'] = df['Text'].str.lower()
```

2.5.3 Removing punctuation

The text might contain extraneous characters such as punctuation marks, special symbols, numbers, etc. which need to be removed as they are not helpful for analysis. I remove them from the dataset by using the following coding, such as `!"#$%&'()*+,-./:;<=>?@[\\]^_`{|}~.`

```
PUNCT_TO_REMOVE = string.punctuation

def remove_punctuation(text):
    """custom function to remove the punctuation"""
    return text.translate(str.maketrans("", "", PUNCT_TO_REMOVE))

df['Text'] = df['Text'].apply(lambda text: remove_punctuation(text))
```

2.5.4 Removing of Stop-words

In this step, I remove stop words in the text, such as "the", "a", "an", etc. These words are not helpful for text mining and thus should be removed.

```

STOPWORDS = set(stopwords.words('english'))

def remove_stopwords(text):
    """custom function to remove the stopwords"""
    return " ".join([word for word in str(text).split() if word not in STOPWORDS])

df['Text'] = df['Text'].apply(lambda text: remove_stopwords(text))

```

2.5.5 Removing URLs

In this step I strip out HTML tags or URLs, as they have no useful information for analysis.

```

def remove_urls(text):
    url_pattern = re.compile(r'https?://\S+|www\.\S+')
    return url_pattern.sub(r'', text)

df['Text'] = df['Text'].apply(lambda text: remove_urls(text))

```

2.6 Checking dataset before next step

```

df.shape

out: (568454, 12)

df.head()

```

Before I build the model, I prefer to check dataset again and ensure that all data are ready to use. Now, all data are cleaned and ready to run into next step.

2.7 Binary classification label

The dataset still contains some textual reviews and their ratings, which is 1 to 5 score with 5 being the best. My goal is to categorize these reviews into positive (4 and 5 score) and negative (1 and 2 score) and remove reviews with a rating 3 because it is unclear whether they are positive or negative. Therefore, I use `y_dict` to convert 1 and 2 in the Score column to 0, and 4 and 5 to 1, indicating that the ratings are divided into two categories, 1 and 2 are negative ratings, and 4 and 5 are positive ratings. Finally, save the transformed Score column into the variable `y`.

```
df['Score'].unique()

df2 = df[df['Score'] != 3]

y_dict = {1:0, 2:0, 4:1, 5:1}

X = df2['Text']

y = df2['Score'].map(y_dict)
```

3. Build Model

3.1 Fitting on train & test set

In previous step (2.4.6), I performed a train-test split on the transformed feature vector dataset. So the data are all ready to fit the model.

I chose Logistics Regression as my predictive model since it is commonly used for binary and multi-class classification tasks. Firstly, I initialized the `LogisticRegression()` class and trained the model using the training dataset.

```
log=LogisticRegression()
```

Then, I evaluated the performance of the trained model on the test dataset, computed the accuracy, and printed it out. Before adjusting the parameters.


```

X_train, X_test, y_train, y_test = train_test_split(X_c, y)

print(' train records: {}'.format(X_train.shape[0]))

ml =log.fit(X_train, y_train)

acc = ml.score(X_test, y_test)

print ('Model Accuracy: {}'.format(acc))

```

I obtained an accuracy of 93.90%, which seems impressive, but there is still room for improvement.

3.2 Fetching Top 20 Positive & Top 20 Negative words

In this step, I output the coefficients of each feature in the trained logistic regression model to find the ranking of the contribution to sentiment classification. Specifically, I used the `get_feature_names()` method in the `CountVectorizer` class to obtain a list of feature names (i.e., words) in the feature vector.

```

w = c.get_feature_names()

```

Then, I obtained the coefficients of each feature from the trained logistic regression model and converted them into a list format. Using the `pandas.DataFrame` class, I converted the dictionary of feature names and coefficients into a data frame for sorting and outputting.

```

coef = ml.coef_.tolist()[0]

coeff_df = pd.DataFrame({'Word' : w, 'Coefficient' : coef})

```

Finally, I sorted the data frame in descending order based on the coefficients and outputted the top 20 positive and negative words with the largest and smallest coefficients, respectively. The former

are typically words related to positive sentiment, while the latter are typically words related to negative sentiment.

```

coeff_df = coeff_df.sort_values(['Coefficient', 'Word'],ascending=False)

print('-Top 20 positive-')

print(coeff_df.head(20).to_string(index=False))

print('\n')

print('-Top 20 negative-')

print(coeff_df.tail(20).to_string(index=False))

```

| -Top 20 positive- | | -Top 20 negative- | |
|-------------------|-------------|-------------------|-------------|
| Word | Coefficient | Word | Coefficient |
| downside | 4.038977 | terriblebr | -2.384423 |
| pleasantly | 3.578166 | ruins | -2.410585 |
| addicting | 2.907093 | overpowers | -2.426812 |
| worries | 2.722353 | ugh | -2.482004 |
| resist | 2.471782 | mediocre | -2.515840 |
| perfectbr | 2.453821 | cancelled | -2.573929 |
| easiest | 2.402521 | disappointment | -2.650258 |
| drawback | 2.390583 | ick | -2.740578 |
| obsessed | 2.349028 | redeeming | -2.746575 |
| skeptical | 2.341225 | dissapointing | -2.773967 |
| hooked | 2.268699 | worst | -2.780524 |
| hesitation | 2.252743 | disappointing | -2.781774 |
| delish | 2.241966 | returnable | -2.785231 |
| solved | 2.237579 | unappealing | -2.837105 |
| trainer | 2.175028 | unacceptable | -2.898754 |
| relaxing | 2.172495 | defeats | -2.981301 |
| dieters | 2.167948 | schar | -2.998371 |
| chocolately | 2.161281 | embarrassed | -3.002553 |
| thankful | 2.154728 | deceptive | -3.205291 |
| tastiest | 2.125783 | undrinkable | -3.566181 |

3.3 Evaluate models by applying multiple NLP techniques & multiple ML algorithms

3.3.1 Text classifier function

```

def text_fit(X, y, nlp_model,ml_model,coef_show=1):

    X_c = nlp_model.fit_transform(X)

    print('features: {}'.format(X_c.shape[1]))

```

```

X_train, X_test, y_train, y_test = train_test_split(X_c, y)

print(' train records: {}'.format(X_train.shape[0]))

print(' test records: {}'.format(X_test.shape[0]))

ml = ml_model.fit(X_train, y_train)

acc = ml.score(X_test, y_test)

print ('Model Accuracy: {}'.format(acc))

if coef_show == 1:

    w = nlp_model.get_feature_names()

    coef = ml.coef_.tolist()[0]

    coeff_df = pd.DataFrame({'Word' : w, 'Coefficient' : coef})

    coeff_df = coeff_df.sort_values(['Coefficient', 'Word'], ascending=[0, 1])

    print('\n')

    print('-Top 20 positive-')

    print(coeff_df.head(20).to_string(index=False))

    print('\n')

    print('-Top 20 negative-')

    print(coeff_df.tail(20).to_string(index=False))

```

In this step, I create a function for implementing the training and evaluation of a text classifier.

This function is –

- ☐ Convert the raw text data X to a numerical feature vector using the `fit_transform()` method of the incoming `nlp_model` object;
- ☐ Divide the data set into training set and test set, using the `train_test_split()` method in `sklearn`;
- ☐ Use the passed in `ml_model` object to fit the training set to get the classifier model `ml`;
- ☐ Predict the test set and calculate the model accuracy `acc`, using the `ml.score()` method;

- If the parameter `coef_show` is 1, the top 20 positive and negative feature coefficients

(Coefficient) after model training are output, that is, the words that have the greatest impact on the classification results, and the output is a pandas data frame.

```
features: 227073
train records: 394360
test records: 131454
Model Accuracy: 0.9387922771463783
```

| -Top 20 positive- | | -Top 20 negative- | |
|-------------------|-------------|-------------------|-------------|
| Word | Coefficient | Word | Coefficient |
| pleasantly | 4.357895 | terriblebr | -2.751530 |
| downside | 3.959523 | schar | -2.797595 |
| addicting | 3.340262 | ripoff | -2.808961 |
| drawback | 3.095615 | allegro | -2.818117 |
| worries | 2.970248 | ruins | -2.867890 |
| easiest | 2.891577 | lacked | -2.885782 |
| delish | 2.820742 | disappointing | -2.895108 |
| obsessed | 2.660533 | lousy | -2.922697 |
| divine | 2.588791 | worst | -2.952930 |
| chocolately | 2.540958 | glowing | -2.974067 |
| soothing | 2.509661 | mediocre | -2.985890 |
| perfectbr | 2.502947 | overpowers | -3.013597 |
| hooked | 2.477601 | dissapointing | -3.080797 |
| excellentbr | 2.475852 | ick | -3.099023 |
| welcome | 2.463927 | embarrassed | -3.176458 |
| settled | 2.435343 | blech | -3.185907 |
| tastiest | 2.422844 | returnable | -3.291776 |
| whim | 2.405144 | unacceptable | -3.555578 |
| skeptical | 2.400208 | undrinkable | -3.908059 |
| trainer | 2.388058 | deceptive | -4.100106 |

Next, I fit the model using the above function. After refinement, I got **93.88%** accuracy.

3.3.2 Prediction function & Confusion Matrix

In this case, I create a prediction function to receive the two models for 'nlp_model' and 'ml_model'. In this case, the function applies a specified NLP model `nlp_model` to transform the text into features and then fits a specified machine learning model `ml_model` to predict `y` based on the features. The function also performs train/test splitting, prints out information about the features and records, computes and prints the confusion matrix and accuracy score of the predictions.

```
def predict(X, y, nlp_model, ml_model):

    X_c = nlp_model.fit_transform(X)
```

```

print('features: {}'.format(X_c.shape[1]))

X_train, X_test, y_train, y_test = train_test_split(X_c, y)

print(' train records: {}'.format(X_train.shape[0]))

print(' test records: {}'.format(X_test.shape[0]))

ml =ml_model.fit(X_train, y_train)

predictions=ml.predict(X_test)

cm=confusion_matrix(predictions,y_test)

print(cm)

acc=accuracy_score(predictions,y_test)

print(acc)

```

```

features: 227073
train records: 394360
test records: 131454
[[ 15003  2486]
 [ 5500 108465]]
0.9392487105755626

```

After I fit the model using the above function, I got **93.92%** accuracy. It is not bad. However, I noticed that some significant coefficients are not meaningful.

3.3.3 Dummy classifier

The reason to use this classifier is that it is to establish a baseline or benchmark for the model's performance, also provides a reference point to compare the actual model's accuracy and decide whether or not the model is doing better than random guessing.

```

text_fit(X, y, c, DummyClassifier(),0)

```

```

features: 227073
train records: 394360
test records: 131454
Model Accuracy: 0.8429184353462048

```

The good news is that I got 84.29% accuracy for this classifier and lower than other twos. Because if the actual learned model does not significantly outperform random guesses, the model needs to be re-evaluated and tuned.

3.3.4 Summary

The third function makes use of DummyClassifier, which is essentially a straightforward random classifier that makes results-only estimates rather than using a genuine learning process. Therefore, the low accuracy of the third function may be due to the result of DummyClassifier.

In contrast, the first and second functions use more complex learning algorithms, so they may be more accurate. In addition, the first function outputs feature coefficients, while the second function outputs stricter evaluation criteria such as confusion matrix and accuracy scores, so these functions provide a more comprehensive model evaluation.

3.4 Logistic regression model on TFIDF vectorizer

In this step, before I perform text classification using a logistic regression model trained on a term frequency-inverse document frequency (TF-IDF) representation of the text data. This helped me fully consider the frequency of each word in the whole text, improving model accuracy and stability. I initialize a TF-IDF vectorizer object with English stop words, then I got the accuracy and output the top 20 positive words and top 20 negative words by using text_fit function. At this step my accuracy is roundly the same as before, but my coefficients have improved. The accuracy is 93.59%.

```
tfidf = TfidfVectorizer(stop_words = 'english')
```

```
text_fit(X, y, tfidf, LogisticRegression())
```

```
features: 227073  
train records: 394360  
test records: 131454  
Model Accuracy: 0.9358558887519589
```

| -Top 20 positive- | | -Top 20 negative- | |
|-------------------|-------------|-------------------|-------------|
| Word | Coefficient | Word | Coefficient |
| great | 14.289092 | yuck | -5.779824 |
| delicious | 12.475956 | hoping | -5.793518 |
| best | 12.102762 | undrinkable | -5.855910 |
| perfect | 11.258012 | waste | -5.999700 |
| excellent | 9.888044 | worse | -6.263415 |
| loves | 9.663801 | bland | -6.428297 |
| highly | 9.218852 | disgusting | -6.434469 |
| love | 8.550390 | stale | -6.440953 |
| wonderful | 8.512197 | weak | -6.696249 |
| amazing | 8.115916 | tasteless | -6.699791 |
| awesome | 7.928247 | return | -7.045196 |
| good | 7.631838 | threw | -7.210885 |
| favorite | 7.226196 | unfortunately | -7.607589 |
| pleased | 7.156822 | horrible | -8.127590 |
| nice | 7.093360 | disappointment | -8.273381 |
| hooked | 6.974443 | disappointed | -8.683583 |
| yummy | 6.959818 | terrible | -9.024415 |
| pleasantly | 6.704298 | awful | -9.068522 |
| smooth | 6.543144 | disappointing | -9.864559 |
| fantastic | 6.430743 | worst | -12.081919 |

```
predict(X, y, tfidf, LogisticRegression())
```

```
features: 227073
train records: 394360
test records: 131454
[[ 14103  2171]
 [  6407 108773]]
0.9347452340742769
```

Then, I got the accuracy and display the Confusion Matrix by prediction function. Accuracy is roughly the same – **93.47%**. However, I noticed that the significant words are now more meaningful and have a higher coefficient magnitude.

3.5 Upvote Prediction

In this step, I am focusing on the score 5 reviews and selecting from its reviews containing 0-20%, 20-40%, 60-80% and 80-100% upvotes and checking the shape of the selected dataset. It contains 154428 data points and 12 features.

```
data = df[df['Score'] == 5]
```

```
data2 = data[data['%upvote'].isin(['0-20%', '20-40%', '60-80%', '80-100%'])]
```

```
data2.shape
```

Then I extract text data from the selected dataset and map upvote to binary output. I'm treating 0-20% and 20-40% as negative category (0), and 60-80% and 80-100% as positive category (1). In the end I came up with a map of upvote tags to a binary output.

```
X = data2['Text']
```

```
y_dict = {'0-20%': 0, '20-40%': 0, '60-80%': 1, '80-100%': 1}
```

```
y = data2['%upvote'].map(y_dict)
```

```
print(y.value_counts())
```

```
1.0    151721
0.0      2707
Name: %upvote, dtype: int64
```

The result of target class 'y' is highly skewed, and I found positive votes are higher than negative votes. This is a case of data imbalance, where using a standard classifier will result in a high error rate,

3.6 Handling imbalanced data using RandomOverSampler

I will re-sample the data to deal with data imbalance and use RandomOverSampler technical. Therefore, I split the data into training and testing sets with 70% of the data used for training and 30% used for testing with the following coding.

```
tf=TfidfVectorizer()
```

```
X_c=tf.fit_transform(X)
```

```
X_train,X_test,y_train,y_test=train_test_split(X_c,y,train_size=0.7)
```

```
y_test.value_counts()
```

Then, I use RandomOverSampler function from imblearn package to oversample the minority class in the training data to deal with class imbalance. In the following coding, I stored the oversampled data in X_train_res and y_train_res.


```
os = RandomOverSampler()

X_train_res, y_train_res = os.fit_sample(X_c, y)
```

Here is the result for original and resampled dataset shapes to show the impact of oversampling. It clearly shows the resampled dataset shape counter for negative votes is 151721.

```
print('Original dataset shape {}'.format(Counter(y)))

print('Resampled dataset shape {}'.format(Counter(y_train_res)))

Original dataset shape Counter({1.0: 151721, 0.0: 2707})
Resampled dataset shape Counter({1.0: 151721, 0.0: 151721})
```

3.7 Tweaking Parameters

Tweaking parameters in this project is the same as finding the best combination of parameters with previous data mining project. I define a logistic regression model with GridSearchCV() function from sklearn to perform parameters tuning on the model and search the highest parameter combinations with the highest accuracy. In this case, I fit the model on the resampled training data using `clf.fit(X_train_res, y_train_res)`, and then predict the target variable on the test data using `y_pred=clf.predict(X_test)`.

```
grid={'C':10.0 **np.arange(-2,3),'penalty':['l1','l2']}

clf=GridSearchCV(estimator=log_class,param_grid=grid,cv=5,n_jobs=-1,scoring='f1_macro')

clf.fit(X_train_res,y_train_res)

y_pred=clf.predict(X_test)

print(confusion_matrix(y_test,y_pred))

print(accuracy_score(y_test,y_pred))
```

The following result is that I got the best combination of parameters with the highest accuracy; getting a score of **98.80%** in the training set, and a score of **99.46%** in the testing set. Here is the best

combination of parameters, confusion matrix, accuracy score in the training set, and accuracy score of the predictions.

```
GridSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,
             param_grid={'C': array([1.e-02, 1.e-01, 1.e+00, 1.e+01, 1.e+02]),
                         'penalty': ['l1', 'l2']},
             scoring='f1_macro')
```

```
In training set, I can get score of : 0.9879728853520786 using {'C': 100.0, 'penalty': 'l2'}
```

```
[[ 813    0]
 [ 249 45267]]
In testing set, I can get the best accuracy is 0.9946253966198277
```

3W

What went well?

Several things went well in this text mining project for sentiment analysis from Amazon reviews.

Firstly, data pre-processing was done efficiently, and the data was cleaned by removing punctuations, URLs, stop words, etc. This helped in reducing the noise in the data and improved the accuracy of the sentiment analysis.

Secondly, the data was explored and visualized using the seaborn library, which provided valuable insights into the data distribution. This helped in identifying patterns in the data and understanding the frequency of words and sentiments in the reviews.

Thirdly, the CountVectorizer function was used to convert the text data into a matrix of token counts, which was then used to train the logistic regression model. The model performed well with an accuracy of 99.41% at the end. The matrix of token counts was fed into the logistic regression model to predict the sentiment of the reviews.

Fourthly, the top positive and negative words were identified using the logistic regression model's coefficients, which helped in understanding the features driving the sentiments.

Finally, the confusion matrix was plotted to evaluate the model's performance. This helped in understanding the model's accuracy and identifying the areas where it was making mistakes.

What did not go well?

One of the major drawbacks of the project was I found the imbalanced data when I predicted upvote. I recognized the target class 'y' is higher skewed and positive votes are much higher than negative votes. I initially overlooked the issue. However, after upon reviewing some pieces of knowledge, I found that imbalanced data can have detrimental effects on the performance of a trained model. This is due to the model being biased towards the majority class, leading to poor performance on the minority class and decreased prediction accuracy overall. Moreover, imbalanced data can result in overfitting and bias in the model. It took me some time to learn how to deal with it and found a suitable way to solve this problem.

Also, my computer's CPU and Jupyter memory limitations may be one of the issues running this project. I often get stuck running Jupyter, forcing me to close everything else and reopen it and run it. Especially when it comes to which classifiers to run.

What I would do differently next time?

I learnt how to handle the imbalanced data by using some techniques. In this case, I used RandomOverSampler function to add the number of negative votes. However, I recognized this method might lead to some overfitting issues and I would try other oversampling techniques next time. Also, next time I will deal with imbalanced data before the model, not at the end.

In addition, when I run this type of code later, I will first clean up the computer memory and close unnecessary background software to ensure that the computer will not be stuck because too many things are running at the same time.

I received a fantastic opportunity to learn and advance my abilities through my work using Python-based text mining. I was able to efficiently alter and get important insights from enormous

volumes of text data because of Python's flexibility and adaptability. It was thrilling to see how different natural language processing tasks, including sentiment analysis and a bag of words, might be carried out using Python tools like NLTK. It was simpler for me to test out various strategies and algorithms because of the abundance of open-source datasets and resources. My curiosity in text mining has been piqued by this encounter, which has inspired me to look into the field's more cutting-edge uses. Overall, I am happy with the project's results and the insightful information discovered.